

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
import os
%matplotlib inline
warnings.filterwarnings('ignore')
```


```
In [2]: df = pd.read_csv("Downloads/CarPrice_Assignment.csv")
```

```
In [3]: df
```

```
Out[3]:
```

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	en
0	1	3	alfa-romero giulia	gas	std	two	convertible	rwd	
1	2	3	alfa-romero stelvio	gas	std	two	convertible	rwd	
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd	
3	4	2	audi 100 ls	gas	std	four	sedan	fwd	
4	5	2	audi 100ls	gas	std	four	sedan	4wd	
...
200	201	-1	volvo 145e (sw)	gas	std	four	sedan	rwd	
201	202	-1	volvo 144ea	gas	turbo	four	sedan	rwd	
202	203	-1	volvo 244dl	gas	std	four	sedan	rwd	
203	204	-1	volvo 246	diesel	turbo	four	sedan	rwd	
204	205	-1	volvo 264gl	gas	turbo	four	sedan	rwd	

205 rows × 26 columns



```
In [4]: #Make a Copy of the Original dataset Which can help me in future
df1 = df.copy(deep=True)
df2 = df.copy(deep=True)
```

Data Preprocessing

In [5]: `df.describe()`

Out[5]:

	car_ID	symboling	wheelbase	carlength	carwidth	carheight	curbweight	eng
count	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205
mean	103.000000	0.834146	98.756585	174.049268	65.907805	53.724878	2555.565854	126
std	59.322565	1.245307	6.021776	12.337289	2.145204	2.443522	520.680204	41
min	1.000000	-2.000000	86.600000	141.100000	60.300000	47.800000	1488.000000	61
25%	52.000000	0.000000	94.500000	166.300000	64.100000	52.000000	2145.000000	97
50%	103.000000	1.000000	97.000000	173.200000	65.500000	54.100000	2414.000000	120
75%	154.000000	2.000000	102.400000	183.100000	66.900000	55.500000	2935.000000	141
max	205.000000	3.000000	120.900000	208.100000	72.300000	59.800000	4066.000000	326

In [6]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_ID                205 non-null   int64
1   symboling              205 non-null   int64
2   CarName                205 non-null   object
3   fueltype               205 non-null   object
4   aspiration              205 non-null   object
5   doornumber             205 non-null   object
6   carbody                205 non-null   object
7   drivewheel             205 non-null   object
8   enginelocation         205 non-null   object
9   wheelbase              205 non-null   float64
10  carlength              205 non-null   float64
11  carwidth                205 non-null   float64
12  carheight              205 non-null   float64
13  curbweight              205 non-null   int64
14  enginetype              205 non-null   object
15  cylindernumber         205 non-null   object
16  enginesize              205 non-null   int64
17  fuelsystem              205 non-null   object
18  boreratio               205 non-null   float64
19  stroke                  205 non-null   float64
20  compressionratio        205 non-null   float64
21  horsepower              205 non-null   int64
22  peakrpm                 205 non-null   int64
23  citympg                 205 non-null   int64
24  highwaympg              205 non-null   int64
25  price                   205 non-null   float64
dtypes: float64(8), int64(8), object(10)
memory usage: 41.8+ KB
```

```
In [7]: df.isnull().sum()
```

```
Out[7]: car_ID          0
        symboling      0
        CarName        0
        fueltype       0
        aspiration     0
        doornumber     0
        carbody        0
        drivewheel     0
        enginelocation 0
        wheelbase      0
        carlength      0
        carwidth       0
        carheight      0
        curbweight     0
        enginetype     0
        cylindernumber 0
        enginesize     0
        fuelsystem     0
        boreratio      0
        stroke         0
        compressionratio 0
        horsepower     0
        peakrpm        0
        citympg        0
        highwaympg     0
        price          0
        dtype: int64
```

```
In [8]: df.columns
```

```
Out[8]: Index(['car_ID', 'symboling', 'CarName', 'fueltype', 'aspiration',
              'doornumber', 'carbody', 'drivewheel', 'enginelocation', 'wheelbase',
              'carlength', 'carwidth', 'carheight', 'curbweight', 'enginetype',
              'cylindernumber', 'enginesize', 'fuelsystem', 'boreratio', 'stroke',
              'compressionratio', 'horsepower', 'peakrpm', 'citympg', 'highwaympg',
              'price'],
              dtype='object')
```

EXPLORATORY DATA ANALYSIS



```
In [9]: #First Of all we seporate categorical and numerical data
        cat = df.select_dtypes(include="object")
        num = df.select_dtypes(include="number")
```



```
In [10]: cat.head()
```



Out[10]:



	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocation	enginetype
0	alfa-romero giulia	gas	std	two	convertible	rwd	front	dol
1	alfa-romero stelvio	gas	std	two	convertible	rwd	front	dol
2	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd	front	ohv
3	audi 100 ls	gas	std	four	sedan	fwd	front	ol
4	audi 100ls	gas	std	four	sedan	4wd	front	ol

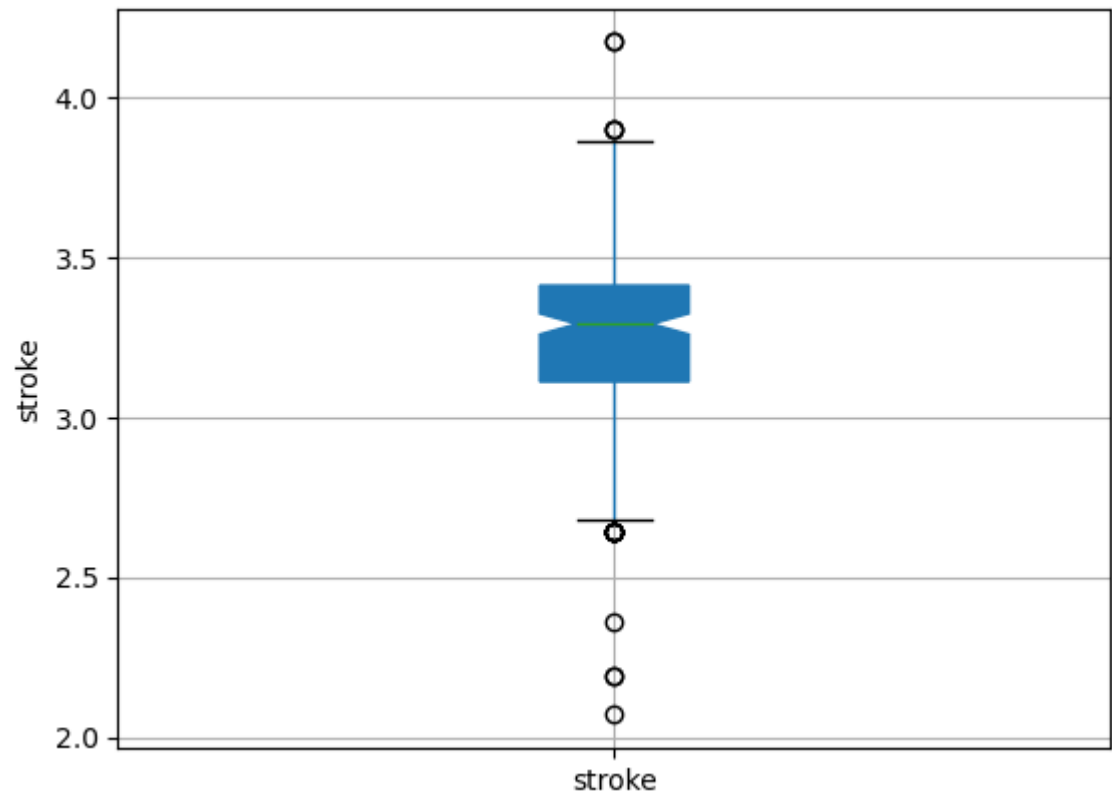

```
In [11]: #for numerical distrubition  
for i in num:  
    num.boxplot(column=i,patch_artist = True, notch = 'True')  
    plt.ylabel(i)  
    plt.show()
```































In [12]: *#we check distrubition of the numerical data*



```
In [13]: for i in num:  
         sns.distplot(df[i], kde = True, color = 'green')  
         plt.show()
```
































```
In [14]: #Now, we can check the categorical features count  
for i in cat:  
    sns.countplot(x=cat[i])  
    plt.show()
```





















```
In [15]: for i in cat:  
          sns.barplot(df, x=df[i], y=df["price"])  
          plt.show()
```











Seperate data in X and Y as well as Split data into train and Test

```
In [16]: df.columns
```

```
Out[16]: Index(['car_ID', 'symboling', 'CarName', 'fueltype', 'aspiration',  
              'doornumber', 'carbody', 'drivewheel', 'enginelocation', 'wheelbase',  
              'carlength', 'carwidth', 'carheight', 'curbweight', 'enginetype',  
              'cylindernumber', 'enginesize', 'fuelsystem', 'boreratio', 'stroke',  
              'compressionratio', 'horsepower', 'peakrpm', 'citympg', 'highwaympg',  
              'price'],  
            dtype='object')
```

```
In [17]: # I am using a df1 data which was copy of the original data set.  
x = df1.drop(["car_ID", "price"], axis=1)  
y = df1["price"]
```


```
In [18]: #for train test split import neccasary library  
from sklearn.model_selection import train_test_split  
train_x, test_x, train_y, test_y = train_test_split(x, y, random_state=50, tes  
t_size=0.2)
```

In [19]: train_x

Out[19]:

	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	engine	location
180	-1	toyota starlet	gas	std	four	sedan	rwd	fror	
38	0	honda civic 1300	gas	std	two	hatchback	fwd	fror	
156	0	toyota mark ii	gas	std	four	sedan	fwd	fror	
40	0	honda accord	gas	std	four	sedan	fwd	fror	
195	-1	volvo 144ea	gas	std	four	wagon	rwd	fror	
...
132	3	saab 99e	gas	std	two	hatchback	fwd	fror	
33	1	honda accord cvcc	gas	std	two	hatchback	fwd	fror	
109	0	peugeot 504 (sw)	gas	std	four	wagon	rwd	fror	
139	2	subaru dl	gas	std	two	hatchback	fwd	fror	
176	-1	toyota corolla	gas	std	four	sedan	fwd	fror	

164 rows × 24 columns



```
In [20]: #we can reset index
train_x.reset_index(inplace=True, drop=True)
test_x.reset_index(inplace=True, drop=True)

train_y.reset_index(inplace=True, drop=True)
test_y.reset_index(inplace=True, drop=True)
```

Encoding using Catboost Encoder

```
In [21]: #for create encoding for input variables we can separate data of numerical and
categorical
train_cat = train_x.select_dtypes(include="object")
train_num = train_x.select_dtypes(include="number")


test_cat = test_x.select_dtypes(include="object")
test_num = test_x.select_dtypes(include="number")
```

In [22]: train_cat

Out[22]:

	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	engine location	enginetype
0	toyota starlet	gas	std	four	sedan	rwd	front	oh
1	honda civic 1300	gas	std	two	hatchback	fwd	front	oh
2	toyota mark ii	gas	std	four	sedan	fwd	front	oh
3	honda accord	gas	std	four	sedan	fwd	front	oh
4	volvo 144ea	gas	std	four	wagon	rwd	front	oh
...
159	saab 99e	gas	std	two	hatchback	fwd	front	oh
160	honda accord cvcc	gas	std	two	hatchback	fwd	front	oh
161	peugeot 504 (sw)	gas	std	four	wagon	rwd	front	oh
162	subaru dl	gas	std	two	hatchback	fwd	front	oh
163	toyota corolla	gas	std	four	sedan	fwd	front	oh

164 rows × 10 columns



In [23]: `import category_encoders as ce`
`encoder = ce.CatBoostEncoder()`
`encoder.fit(train_cat, train_y)`

Out[23]:

```

CatBoostEncoder
CatBoostEncoder(cols=['CarName', 'fueltype', 'aspiration', 'doornumber',
                        'carbody', 'drivewheel', 'engine location', 'enginetype',
                        'cylindernumber', 'fuelsystem'])

```

In [24]: `train_cat = encoder.transform(train_cat)`
`test_cat = encoder.transform(test_cat)`

In [25]: train_cat

Out[25]:

	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	engine
0	13121.513551	13434.368051	13069.674208	13789.064262	14623.968470	20484.766260	1329
1	13685.540652	13434.368051	13069.674208	13550.384829	10777.771565	9273.942533	1329
2	10282.385163	13434.368051	13069.674208	13789.064262	14623.968470	9273.942533	1329
3	10625.180217	13434.368051	13069.674208	13789.064262	14623.968470	9273.942533	1329
4	15381.846884	13434.368051	13069.674208	13789.064262	12587.501767	20484.766260	1329
...
159	14718.513551	13434.368051	13069.674208	13550.384829	10777.771565	9273.942533	1329
160	13685.540652	13434.368051	13069.674208	13550.384829	10777.771565	9273.942533	1329
161	13685.540652	13434.368051	13069.674208	13789.064262	12587.501767	20484.766260	1329
162	9609.708130	13434.368051	13069.674208	13550.384829	10777.771565	9273.942533	1329
163	11270.135163	13434.368051	13069.674208	13789.064262	14623.968470	9273.942533	1329

164 rows × 10 columns

In [26]: *# Now, we concat the both categorical and numerical data*
train_x1 = pd.concat([train_num, train_cat], axis=1)
test_x1 = pd.concat([test_num, test_cat], axis=1)

In [27]: train_x1.head()

Out[27]:

	symboling	wheelbase	carlength	carwidth	carheight	curbweight	enginesize	boreratio	strok
0	-1	104.5	187.8	66.5	54.1	3131	171	3.27	3.3
1	0	96.5	167.5	65.2	53.3	2289	110	3.15	3.5
2	0	95.7	166.3	64.4	53.0	2081	98	3.19	3.0
3	0	96.5	175.4	62.5	54.1	2372	110	3.15	3.5
4	-1	104.3	188.8	67.2	57.5	3034	141	3.78	3.1

5 rows × 24 columns

Scaling Using MinmaxScaler

```
In [28]: from sklearn.preprocessing import MinMaxScaler, StandardScaler, RobustScaler
scaler = MinMaxScaler()
scaler.fit(train_x1)
```

```
Out[28]: ▼ MinMaxScaler
MinMaxScaler()
```


```
In [29]: train_x1 = pd.DataFrame(scaler.transform(train_x1), columns=train_x1.columns)
test_x1 = pd.DataFrame(scaler.transform(test_x1), columns=test_x1.columns)
```

```
In [30]: train_x1.head()
```

```
Out[30]:
```

	symboling	wheelbase	carlength	carwidth	carheight	curbweight	enginesize	boreratio	str
0	0.2	0.617241	0.759350	0.516667	0.525000	0.637316	0.415094	0.521429	0.609
1	0.4	0.341379	0.429268	0.408333	0.458333	0.310706	0.184906	0.435714	0.719
2	0.4	0.313793	0.409756	0.341667	0.433333	0.230023	0.139623	0.464286	0.457
3	0.4	0.341379	0.557724	0.183333	0.525000	0.342901	0.184906	0.435714	0.719
4	0.2	0.610345	0.775610	0.575000	0.808333	0.599690	0.301887	0.885714	0.514

5 rows × 24 columns



Model Building And Evaluation

```
In [31]: from sklearn.linear_model import LinearRegression
from sklearn.neighbors import KNeighborsRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, AdaBoostRegressor
from sklearn.svm import SVR
import xgboost as Xgb
import catboost as cb
```

```
In [32]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

```
In [33]: #LinerRegression
log_model = LinearRegression()
log_model.fit(train_x1, train_y)
pred_log = log_model.predict(test_x1)
mea_log = mean_absolute_error(test_y, pred_log)
mea_log
```

```
Out[33]: 1862.0754026991408
```

```
In [34]: log_model.score(train_x1, train_y)
```

```
Out[34]: 0.9145827618690593
```

```
In [35]: log_model.score(test_x1, test_y)
```

```
Out[35]: 0.8439860393076992
```

```
In [36]: #KNEARASTNEIGHBORS REGRESSOR
knn_model = KNeighborsRegressor(n_neighbors=10)
knn_model.fit(train_x1, train_y)
pred_knn = knn_model.predict(test_x1)
mea_knn = mean_absolute_error(test_y, pred_knn)
mea_knn
```

```
Out[36]: 2346.491056097561
```

```
In [37]: knn_model.score(train_x1, train_y)
```

```
Out[37]: 0.7583043733079329
```

```
In [38]: knn_model.score(test_x1, test_y)
```

```
Out[38]: 0.7295957720460675
```

```
In [39]: # SUPPORT VECTOR REGRESSOR
svm_model = SVR(kernel="rbf")
svm_model.fit(train_x1, train_y)
pred_svm = svm_model.predict(test_x1)
mea_svm = mean_absolute_error(test_y, pred_svm)
mea_svm
```

```
Out[39]: 4405.829720847327
```

```
In [40]: #DECISION TREE REGRESSOR
dt_model = DecisionTreeRegressor(random_state=50)
dt_model.fit(train_x1, train_y)
pred_dt = dt_model.predict(test_x1)
mea_dt = mean_absolute_error(test_y, pred_dt)
mea_dt
```

```
Out[40]: 1873.0365853658536
```

```
In [41]: dt_model.score(train_x1, train_y)
```

```
Out[41]: 0.9990939296264457
```

```
In [42]: dt_model.score(test_x1, test_y)
```

```
Out[42]: 0.835366233138423
```



```
In [43]: #XGBOOST REGRESSOR
xgb_model = Xgb.XGBRegressor(n_estimators=100)
xgb_model.fit(train_x1, train_y)
pred_xgb = xgb_model.predict(test_x1)
mea_xgb = mean_absolute_error(test_y, pred_xgb)
mea_xgb
```

Out[43]: 1492.0071217606708

```
In [44]: xgb_model.score(train_x1, train_y)
```

Out[44]: 0.9990933614444724

```
In [45]: xgb_model.score(test_x1, test_y)
```

Out[45]: 0.904477001109033

```
In [46]: #RANDOM FOREST REGRESSOR
rfc_model = RandomForestRegressor(random_state=50)
rfc_model.fit(train_x1, train_y)
pred_rfc = rfc_model.predict(test_x1)
mea_rfc = mean_absolute_error(test_y, pred_rfc)
mea_rfc
```

Out[46]: 1312.5139427642275

```
In [47]: rfc_model.score(train_x1, train_y)
```

Out[47]: 0.9891398153002534

```
In [48]: rfc_model.score(test_x1, test_y)
```

Out[48]: 0.9235170839824554

```
In [49]: #ADABOOST REGRESSOR
from sklearn.ensemble import AdaBoostRegressor
adb_model = AdaBoostRegressor(random_state=50)
adb_model.fit(train_x1, train_y)
pred_adb = adb_model.predict(test_x1)
mea_adb = mean_absolute_error(test_y, pred_adb)
mea_adb
```

Out[49]: 1722.7102183490472

```
In [50]: adb_model.score(train_x1, train_y)
```

Out[50]: 0.960733954988325

```
In [51]: adb_model.score(test_x1, test_y)
```

Out[51]: 0.8941449598689841

HYPERPARAMETER TUNING

In [52]: `from sklearn.model_selection import GridSearchCV`

In [53]: `#HYPERPARAMETER TUNING OF KNN`
`knn = KNeighborsRegressor()`
`params_knn = {'algorithm': ['auto', 'ball_tree', 'kd_tree', 'brute'], 'weights': ['uniform', 'distance'],`
`'n_neighbors': [1, 25, 14, 13, 26, 85, 45]}`
`clf2 = GridSearchCV(knn, params_knn, cv=5, scoring="neg_mean_absolute_error")`
`clf2.fit(train_x1, train_y)`
`print(clf2.best_params_)`
`print(-(clf2.best_score_))`

`{'algorithm': 'auto', 'n_neighbors': 14, 'weights': 'distance'}`
`2494.429238083901`

In [54]: `#HYPERPARAMETER TUNING OF SUPPORT VECTOR`
`svm = SVR()`
`params_svm = {"gamma": ["scale", "auto"]}`
`clf4 = GridSearchCV(svm, params_svm, cv=5, scoring="neg_mean_absolute_error")`
`clf4.fit(train_x1, train_y)`
`print(clf4.best_params_)`
`print(-(clf4.best_score_))`

`{'gamma': 'scale'}`
`5857.0948700388735`

In [55]: `#HYPERPARAMETER TUNING OF DECISION TREE`
`dt = DecisionTreeRegressor()`
`params_dt = { 'max_depth': [1, 25, 14, 13, 45, 75, 26], 'splitter': ['best', 'random'] }`
`clf5 = GridSearchCV(dt, params_dt, cv=5, scoring="neg_mean_absolute_error")`
`clf5.fit(train_x1, train_y)`
`print(clf5.best_params_)`
`print(-(clf5.best_score_))`

`{'max_depth': 26, 'splitter': 'best'}`
`2079.834215151515`

In [56]: `#HYPERPARAMETER TUNING OF RANDOMFOREST`
`rfc = RandomForestRegressor()`
`params_rfc = {"n_estimators": [10, 15, 125, 10, 8, 85], "max_depth": [10, 25, 48, 85, 42, 3]}`
`clf6 = GridSearchCV(rfc, params_rfc, cv=5, scoring="neg_mean_absolute_error")`
`clf6.fit(train_x1, train_y)`
`print(clf6.best_params_)`
`print(-(clf6.best_score_))`

`{'max_depth': 10, 'n_estimators': 10}`
`1608.045970352032`

```
In [57]: #HYPERPERAMETER TUNING OF XGBOOST
xgb = Xgb.XGBRegressor()
params_xgb = {'eta': [0.1, 0.2, 0.3,0.4,0.5], 'n_estimators' : [10, 50, 100,12,15], 'max_depth': [3, 6, 9,14]}
clf7 = GridSearchCV(xgb, params_xgb, cv=5, scoring="neg_mean_absolute_error")
clf7.fit(train_x1, train_y)
print(clf7.best_params_)
print(-(clf7.best_score_))
```

```
{'eta': 0.3, 'max_depth': 3, 'n_estimators': 12}
1654.4771265676786
```

```
In [58]: #HYPERPERAMETER TUNING OF ADABOOST
adb = AdaBoostRegressor()
params_adb = {'n_estimators' : [10, 50, 100,12,15]}
clf8 = GridSearchCV(adb, params_adb, cv=5, scoring="neg_mean_absolute_error")
clf8.fit(train_x1, train_y)
print(clf8.best_params_)
print(-(clf8.best_score_))
```

```
{'n_estimators': 12}
1942.4232639553607
```

```
In [59]: #best parameter for model
print("KNeighborsRegressor score is :", clf2.best_params_)
print("Support vector machine score is :", clf4.best_params_)
print("DecisionTreeRegressor score is :", clf5.best_params_)
print("RandomForestRegressor score is :", clf6.best_params_)
print("XGBOOST score is :", clf7.best_params_)
print("AdaBoostRegressor score is :", clf8.best_params_)
```

```
KNeighborsRegressor score is : {'algorithm': 'auto', 'n_neighbors': 14, 'weights': 'distance'}
Support vector machine score is : {'gamma': 'scale'}
DecisionTreeRegressor score is : {'max_depth': 26, 'splitter': 'best'}
RandomForestRegressor score is : {'max_depth': 10, 'n_estimators': 10}
XGBOOST score is : {'eta': 0.3, 'max_depth': 3, 'n_estimators': 12}
AdaBoostRegressor score is : {'n_estimators': 12}
```

```
In [60]: #Score for all model
print("KNeighborsRegressor score is :", -clf2.best_score_)
print("Support vector machine score is :", -clf4.best_score_)
print("DecisionTreeRegressor score is :", -clf5.best_score_)
print("RandomForestRegressor score is :", -clf6.best_score_)
print("XGBOOST score is :", -clf7.best_score_)
print("AdaBoostRegressor score is :", -clf8.best_score_)
```

```
KNeighborsRegressor score is : 2494.429238083901
Support vector machine score is : 5857.0948700388735
DecisionTreeRegressor score is : 2079.834215151515
RandomForestRegressor score is : 1608.045970352032
XGBOOST score is : 1654.4771265676786
AdaBoostRegressor score is : 1942.4232639553607
```

Feature Selection

```
In [61]: corr = train_x1.corr()
corr.style.background_gradient(cmap='coolwarm')
```

Out[61]:

	symboling	wheelbase	carlength	carwidth	carheight	curbweight	enginesize
symboling	1.000000	-0.578304	-0.390113	-0.248365	-0.526122	-0.250262	-0.134253
wheelbase	-0.578304	1.000000	0.867582	0.771214	0.591112	0.759793	0.529871
carlength	-0.390113	0.867582	1.000000	0.827044	0.470846	0.876853	0.674825
carwidth	-0.248365	0.771214	0.827044	1.000000	0.237273	0.867648	0.735210
carheight	-0.526122	0.591112	0.470846	0.237273	1.000000	0.242780	0.031296
curbweight	-0.250262	0.759793	0.876853	0.867648	0.242780	1.000000	0.863212
enginesize	-0.134253	0.529871	0.674825	0.735210	0.031296	0.863212	1.000000
boreratio	-0.162739	0.472559	0.602451	0.558945	0.153063	0.651217	0.567905
stroke	-0.094843	0.233869	0.159160	0.187711	0.021950	0.177018	0.175712
compressionratio	-0.138442	0.276306	0.162644	0.175197	0.256479	0.156818	0.030569
horsepower	0.041720	0.320702	0.542183	0.639089	-0.132944	0.760007	0.812254
peakrpm	0.298979	-0.359770	-0.263725	-0.199603	-0.323789	-0.230956	-0.208542
citympg	-0.010920	-0.444372	-0.667301	-0.639629	-0.015309	-0.757775	-0.650781
highwaympg	0.071724	-0.524224	-0.706315	-0.670228	-0.079328	-0.800212	-0.673568
CarName	-0.010792	0.258350	0.307601	0.369344	0.068102	0.382338	0.339512
fueltype	-0.168656	0.340227	0.218323	0.224837	0.294276	0.218850	0.064077
aspiration	-0.139669	0.309735	0.256531	0.291672	0.155802	0.329624	0.094739
doornumber	-0.660961	0.485944	0.419589	0.202060	0.555275	0.210491	0.039682
carbody	0.059908	0.079006	0.211546	0.132693	0.032433	0.237191	0.382170
drivewheel	-0.080845	0.470344	0.533570	0.515855	-0.001609	0.691589	0.563847
enginelocation	0.237181	-0.212719	-0.059901	-0.061146	-0.124708	0.052068	0.213736
enginetype	0.067769	0.154616	0.304882	0.349761	-0.136303	0.502473	0.585707
cylindernumber	-0.007060	0.314598	0.428817	0.571120	-0.051041	0.634563	0.745415
fuelsystem	0.015097	0.472859	0.633303	0.602246	0.135930	0.674265	0.544771

```
In [62]: def correlation(dataset, threshold):
col_corr = set()
corr_matrix = dataset.corr()
for i in range(len(corr_matrix.columns)):
    for j in range(i):
        if abs(corr_matrix.iloc[i,j]) > threshold:
            colname = corr_matrix.columns[i]
            col_corr.add(colname)
return col_corr
```

```
In [63]: corr_features = correlation(train_x1, 0.7)
len(set(corr_features))
```

Out[63]: 9

```
In [64]: corr_features
```

```
Out[64]: {'carlength',
'carwidth',
'citympg',
'curbweight',
'cylindernumber',
'enginesize',
'fueltype',
'highwaympg',
'horsepower'}
```

```
In [65]: from sklearn.preprocessing import LabelEncoder
from sklearn.feature_selection import SelectKBest, chi2
label_encoder = LabelEncoder()
categorical_labels = label_encoder.fit_transform(train_y)

selector = SelectKBest(score_func=chi2, k=9)
selector.fit(train_x1, categorical_labels)
```

```
Out[65]: SelectKBest
SelectKBest(k=9, score_func=<function chi2 at 0x0000026D07CE2320>)
```

```
In [66]: train_x1.columns[selector.get_support()]
```

```
Out[66]: Index(['compressionratio', 'fueltype', 'aspiration', 'doornumber',
'drivewheel', 'engine location', 'enginetype', 'cylindernumber',
'fuelsystem'],
dtype='object')
```


```
In [67]: from sklearn.ensemble import ExtraTreesClassifier
model = ExtraTreesClassifier()
model.fit(train_x1, categorical_labels)
```

```
Out[67]: ExtraTreesClassifier
ExtraTreesClassifier()
```

```
In [68]: print(model.feature_importances_)
```

```
[0.03838191 0.0470145  0.04979188 0.04186319 0.06225254 0.12788082  
 0.03876581 0.03906442 0.03808178 0.03694287 0.04472997 0.04336357  
 0.05476364 0.05119548 0.09452813 0.00513165 0.01301257 0.03345145  
 0.05890118 0.01966169 0.00134083 0.02081258 0.01515469 0.02391286]
```

```
In [69]: feat_importance = pd.Series(model.feature_importances_, index=x.columns)  
feat_importance.nlargest(26).plot(kind="barh")  
plt.show()
```



```
In [70]: #Apply SelectKbest class to extract top Features
from sklearn.feature_selection import SelectKBest, chi2
bestfeatures = SelectKBest(score_func=chi2, k=7)
fit = bestfeatures.fit(train_x1,categorical_labels)
dfscores = pd.DataFrame(fit.scores_)
dfcolumns = pd.DataFrame(x.columns)
features = pd.concat([dfcolumns, dfscores], axis=1)
features.columns = ["specs", "score"]
features
```

Out[70]:

	specs	score
0	symboling	18.022609
1	CarName	16.372104
2	fueltype	12.133704
3	aspiration	10.996294
4	doornumber	13.152871
5	carbody	16.141302
6	drivewheel	16.246778
7	enginelocation	10.689274
8	wheelbase	5.988593
9	carlength	46.766794
10	carwidth	19.838707
11	carheight	15.116577
12	curbweight	16.580375
13	enginetype	14.208251
14	cylindernumber	14.744994
15	enginesize	143.533333
16	fuelsystem	137.000000
17	boreratio	69.236559
18	stroke	32.983696
19	compressionratio	95.168704
20	horsepower	161.000000
21	peakrpm	90.965309
22	citympg	80.619337
23	highwaympg	57.424042

```
In [71]: from sklearn.ensemble import RandomForestClassifier
fe_model = RandomForestClassifier(random_state=50)

fe_model.fit(train_x1, categorical_labels)
```

```
Out[71]: ▼      RandomForestClassifier
RandomForestClassifier(random_state=50)
```

```
In [72]: feature_scores = pd.Series(fe_model.feature_importances_, index=train_x.columns)
         .sort_values(ascending=False)
         feature_scores
```

```
Out[72]: carbody          0.110586
cylindernumber    0.080961
doornumber        0.072819
carwidth          0.056263
curbweight        0.054876
fueltype          0.054374
aspiration        0.052511
enginetype        0.050929
carlength         0.049689
drivewheel        0.048289
CarName           0.045783
stroke            0.045180
enginelocation    0.044269
carheight         0.043742
wheelbase         0.043452
symboling         0.039086
boreratio         0.028178
peakrpm           0.019195
highwaympg        0.017722
compressionratio  0.013915
citympg           0.011500
fuelsystem        0.010661
enginesize        0.003963
horsepower        0.002058
dtype: float64
```

After Feature Selection drop the features of noises and regenerate the model

```
In [73]: df2.columns
```

```
Out[73]: Index(['car_ID', 'symboling', 'CarName', 'fueltype', 'aspiration',
               'doornumber', 'carbody', 'drivewheel', 'enginelocation', 'wheelbase',
               'carlength', 'carwidth', 'carheight', 'curbweight', 'enginetype',
               'cylindernumber', 'enginesize', 'fuelsystem', 'boreratio', 'stroke',
               'compressionratio', 'horsepower', 'peakrpm', 'citympg', 'highwaympg',
               'price'],
              dtype='object')
```



```
In [74]: X = df2.drop(["car_ID", "enginesize", "horsepower", "price", "enginelocation",
'carlength', 'carwidth', 'carheight', 'curbweight', 'fuelsystem', 'citympg'],
axis=1)
Y = df2["price"]
```

```
In [75]: X
```

```
Out[75]:
```

	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	wheelbase
0	3	alfa-romero giulia	gas	std	two	convertible	rwd	88.6
1	3	alfa-romero stelvio	gas	std	two	convertible	rwd	88.6
2	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd	94.5
3	2	audi 100 ls	gas	std	four	sedan	fwd	99.8
4	2	audi 100ls	gas	std	four	sedan	4wd	99.4
...
200	-1	volvo 145e (sw)	gas	std	four	sedan	rwd	109.1
201	-1	volvo 144ea	gas	turbo	four	sedan	rwd	109.1
202	-1	volvo 244dl	gas	std	four	sedan	rwd	109.1
203	-1	volvo 246	diesel	turbo	four	sedan	rwd	109.1
204	-1	volvo 264gl	gas	turbo	four	sedan	rwd	109.1

205 rows × 15 columns

```
In [76]: train_X, test_X, train_Y, test_Y = train_test_split(X,Y, random_state=50, test
_size=0.2)
```

```
In [77]: train_X.head()
```

```
Out[77]:
```

	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	wheelbase	e
180	-1	toyota starlet	gas	std	four	sedan	rwd	104.5	
38	0	honda civic 1300	gas	std	two	hatchback	fwd	96.5	
156	0	toyota mark ii	gas	std	four	sedan	fwd	95.7	
40	0	honda accord	gas	std	four	sedan	fwd	96.5	
195	-1	volvo 144ea	gas	std	four	wagon	rwd	104.3	

```
In [78]: #we can reset index
train_X.reset_index(inplace=True, drop=True)
test_X.reset_index(inplace=True, drop=True)

train_Y.reset_index(inplace=True, drop=True)
test_Y.reset_index(inplace=True, drop=True)
```

```
In [79]: train_cat1 = train_X.select_dtypes(include="object")
train_num1 = train_X.select_dtypes(include="number")

test_cat1 = test_X.select_dtypes(include="object")
test_num1 = test_X.select_dtypes(include="number")
```

```
In [80]: train_cat1
```

Out[80]:

	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	enginetype	cylindernumt
0	toyota starlet	gas	std	four	sedan	rwd	dohc	
1	honda civic 1300	gas	std	two	hatchback	fwd	ohc	fc
2	toyota mark ii	gas	std	four	sedan	fwd	ohc	fc
3	honda accord	gas	std	four	sedan	fwd	ohc	fc
4	volvo 144ea	gas	std	four	wagon	rwd	ohc	fc
...	
159	saab 99e	gas	std	two	hatchback	fwd	ohc	fc
160	honda accord cvcc	gas	std	two	hatchback	fwd	ohc	fc
161	peugeot 504 (sw)	gas	std	four	wagon	rwd	l	fc
162	subaru dl	gas	std	two	hatchback	fwd	ohcf	fc
163	toyota corolla	gas	std	four	sedan	fwd	ohc	fc

164 rows × 8 columns



```
In [81]: encoder.fit(train_cat1, train_Y)
train_cat1 = encoder.transform(train_cat1)
test_cat1 = encoder.transform(test_cat1)
```

In [82]: train_cat1

Out[82]:

	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	eng
0	13121.513551	13434.368051	13069.674208	13789.064262	14623.968470	20484.766260	18707
1	13685.540652	13434.368051	13069.674208	13550.384829	10777.771565	9273.942533	11808
2	10282.385163	13434.368051	13069.674208	13789.064262	14623.968470	9273.942533	11808
3	10625.180217	13434.368051	13069.674208	13789.064262	14623.968470	9273.942533	11808
4	15381.846884	13434.368051	13069.674208	13789.064262	12587.501767	20484.766260	11808
...
159	14718.513551	13434.368051	13069.674208	13550.384829	10777.771565	9273.942533	11808
160	13685.540652	13434.368051	13069.674208	13550.384829	10777.771565	9273.942533	11808
161	13685.540652	13434.368051	13069.674208	13789.064262	12587.501767	20484.766260	14742
162	9609.708130	13434.368051	13069.674208	13550.384829	10777.771565	9273.942533	15047
163	11270.135163	13434.368051	13069.674208	13789.064262	14623.968470	9273.942533	11808

164 rows × 8 columns

In [83]: train_X = pd.concat([train_num1, train_cat1], axis=1)
test_X = pd.concat([test_num1, test_cat1], axis=1)

In [84]: train_X

Out[84]:

	symboling	wheelbase	boreratio	stroke	compressionratio	peakrpm	highwaympg	CarName
0	-1	104.5	3.27	3.35	9.20	5200	24	13121.51
1	0	96.5	3.15	3.58	9.00	5800	33	13685.54
2	0	95.7	3.19	3.03	9.00	4800	37	10282.38
3	0	96.5	3.15	3.58	9.00	5800	33	10625.18
4	-1	104.3	3.78	3.15	9.50	5400	28	15381.84
...
159	3	99.1	3.54	3.07	9.31	5250	28	14718.51
160	1	93.7	2.91	3.41	9.20	6000	34	13685.54
161	0	114.2	3.46	3.19	8.40	5000	24	13685.54
162	2	93.7	3.62	2.64	8.70	4400	31	9609.70
163	-1	102.4	3.31	3.54	8.70	4200	32	11270.13

164 rows × 15 columns

```
In [85]: scaler.fit(train_X)
train_X = pd.DataFrame(scaler.transform(train_X), columns=train_X.columns)
test_X = pd.DataFrame(scaler.transform(test_X), columns=test_X.columns)
```

```
In [86]: test_X.head()
```

```
Out[86]:
```

	symboling	wheelbase	boreratio	stroke	compressionratio	peakrpm	highwaympg	CarName
0	0.4	0.951724	0.657143	0.057143	0.08750	0.346939	0.210526	0.372065
1	0.8	0.406897	0.771429	0.680952	0.14375	0.265306	0.368421	0.264484
2	0.4	0.503448	0.550000	0.533333	0.12500	0.040816	0.315789	0.264484
3	0.4	1.182759	0.900000	0.609524	0.06250	0.142857	0.000000	0.264484
4	0.4	0.358621	0.771429	0.271429	0.04375	0.265306	0.342105	0.264484

```
In [87]: #KNeighborsRegressor
knn1 = KNeighborsRegressor(algorithm="auto", n_neighbors=14, weights="distance")
knn1.fit(train_X, train_Y)
pred1 = knn1.predict(test_X)
mea_knn1 = mean_absolute_error(test_Y, pred1)
mea_knn1
```

```
Out[87]: 2336.4669472990554
```

```
In [88]: knn1.score(test_X, test_Y)
```

```
Out[88]: 0.667114714416121
```

```
In [89]: #SVR
svm1 = SVR(gamma="scale")
svm1.fit(train_X, train_Y)
pred2 = svm1.predict(test_X)
mea_svm = mean_absolute_error(test_Y, pred2)
mea_svm
```

```
Out[89]: 4411.585208835385
```

```
In [90]: #DecisionTreeRegressor
dt1 = DecisionTreeRegressor( max_depth=45, splitter="best")
dt1.fit(train_X, train_Y)
pred3 = dt1.predict(test_X)
mea_dt = mean_absolute_error(test_Y, pred3)
mea_dt
```

```
Out[90]: 1500.621951219512
```

```
In [91]: dt1.score(test_X, test_Y)
```

```
Out[91]: 0.8344006282029793
```

```
In [92]: #RandomForestRegressor
rfc1 = RandomForestRegressor(max_depth=25 ,n_estimators= 15)
rfc1.fit(train_X, train_Y)
pred4 = rfc1.predict(test_X)
mea_rfc = mean_absolute_error(test_Y, pred4)
mea_rfc
```

Out[92]: 1178.684498102981

```
In [93]: rfc1.score(test_X, test_Y)
```

Out[93]: 0.942488779940313

```
In [94]: rfc1.score(train_X, train_Y)
```

Out[94]: 0.9789761134166135

```
In [95]: #XGBRegressor
xgb = Xgb.XGBRegressor(eta=0.3 ,max_depth=3 ,n_estimators= 12)
xgb.fit(train_X, train_Y)
pred5 = xgb.predict(test_X)
mea_xgb = mean_absolute_error(test_Y, pred5)
mea_xgb
```

Out[95]: 1363.2071265243903

```
In [96]: xgb.score(train_X, train_Y)
```

Out[96]: 0.9569576336137279

```
In [97]: xgb.score(test_X, test_Y)
```

Out[97]: 0.9326547476323311

```
In [98]: #AdaBoostRegressor
adb = AdaBoostRegressor(n_estimators= 15)
adb.fit(train_X, train_Y)
pred6 = adb.predict(test_X)
mea_adb1 = mean_absolute_error(test_Y, pred6)
mea_adb1
```

Out[98]: 1418.285329668117

```
In [99]: adb.score(test_X, test_Y)
```

Out[99]: 0.9243458301093536

```
In [100]: print('KNeighborsRegressor MAE is: ', mea_knn1)
print('KNeighborsRegressor training score is: ', knn1.score(train_X, train_Y))
print('KNeighborsRegressor test score is : ', knn1.score(test_X, test_Y))
print()
print('DecisionTreeRegressor MAE is :', mea_dt)
print('DecisionTreeRegressor training score is: ', dt1.score(train_X, train_Y))
print('DecisionTreeRegressor test score is : ', dt1.score(test_X, test_Y))
print()
print('RandomForestRegressor MAE is :', mea_rfc)
print('RandomForestRegressor training score is: ', rfc1.score(train_X, train_Y))
print('RandomForestRegressor test score is : ', rfc1.score(test_X, test_Y))
print()
print('xgboostregression MAE is :', mea_xgb)
print('xgboostregression training score is: ', xgb.score(train_X, train_Y))
print('xgboostregression test score is : ', xgb.score(test_X, test_Y))
print()
print('AdaBoostRegressor MAE is :', mea_adb1)
print('AdaBoostRegressor training score is: ', adb.score(train_X, train_Y))
print('AdaBoostRegressor test score is : ', adb.score(test_X, test_Y))
print()
```

KNeighborsRegressor MAE is: 2336.4669472990554
KNeighborsRegressor training score is: 0.9989318339613448
KNeighborsRegressor test score is : 0.667114714416121

DecisionTreeRegressor MAE is : 1500.621951219512
DecisionTreeRegressor training score is: 0.9989318339613448
DecisionTreeRegressor test score is : 0.8344006282029793

RandomForestRegressor MAE is : 1178.684498102981
RandomForestRegressor training score is: 0.9789761134166135
RandomForestRegressor test score is : 0.942488779940313

xgboostregression MAE is : 1363.2071265243903
xgboostregression training score is: 0.9569576336137279
xgboostregression test score is : 0.9326547476323311

AdaBoostRegressor MAE is : 1418.285329668117
AdaBoostRegressor training score is: 0.9440312926868266
AdaBoostRegressor test score is : 0.9243458301093536

Testing the Model

In [101]: X.head(1)

Out[101]:

	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	wheelbase	eng
0	3	alfa-romero giulia	gas	std	two	convertible	rwd	88.6	

```
In [102]: new_df = {'symboling': 0, 'CarName': "bmw X2", 'fueltype' : 'gas', 'aspiration': "turbo",
                  "doornumber": "four", "carbody": "convertible", 'drivewheel': "rwd",
                  'wheelbase': 98.0,
                  'enginetype': 'ohcf', 'cylindernumber': 'four', 'boreratio': 4.51, 'stroke': 3.68,
                  'compressionratio': 9.0, 'peakrpm': 5500, 'highwaympg': 22}
index = [0]
```

In [103]: new_df = pd.DataFrame(new_df, index=index)

In [104]: new_df

Out[104]:

	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	wheelbase	eng
0	0	bmw X2	gas	turbo	four	convertible	rwd	98.0	

```
In [105]: new_cat = new_df.select_dtypes(include="object")
new_num = new_df.select_dtypes(include="number")
```

In [106]: new_cat = encoder.transform(new_cat)

In [107]: new_cat

Out[107]:

	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	enginetype
0	13685.540652	13434.368051	16720.882416	13789.064262	20718.36295	20484.76626	15047.1185

In [108]: new_df = pd.concat([new_num, new_cat], axis=1)

In [109]: new_df

Out[109]:

	symboling	wheelbase	boreratio	stroke	compressionratio	peakrpm	highwaympg	CarName
0	0	98.0	4.51	3.68	9.0	5500	22	13685.5406

```
In [110]: new_df =pd.DataFrame(scaler.transform(new_df), columns=new_df.columns)
new_df
```

```
Out[110]:
```

	symboling	wheelbase	boreratio	stroke	compressionratio	peakrpm	highwaympg	CarName
0	0.4	0.393103	1.407143	0.766667	0.125	0.55102	0.157895	0.264484

```
In [111]: Pred_new = rfc1.predict(new_df)
```

```
In [112]: Pred_new
```

```
Out[112]: array([15973.51111111])
```

```
In [113]: X["enginetype"].unique()
```

```
Out[113]: array(['dohc', 'ohcv', 'ohc', 'l', 'rotor', 'ohcf', 'dohcv'], dtype=object)
```

```
In [114]: new_df1 = {'symboling': 3, 'CarName':"wagon R", 'fueltype' : 'gas', 'aspiration':"std" ,
                    "doornumber":"four", "carbody":"wagon", 'drivewheel': "rwd", 'wheelbase':98.0,
                    'enginetype':'dohc', 'cylindernumber':'three', 'boreratio': 4.51,
                    'stroke':4,
                    'compressionratio':8.0, 'peakrpm': 5000, 'highwaympg': 29}
index = [0]
```

```
In [115]: new_df1 = pd.DataFrame(new_df1, index=index)
```

```
In [116]: new_df1
```

```
Out[116]:
```

	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	wheelbase	engine
0	3	wagon R	gas	std	four	wagon	rwd	98.0	

```
In [117]: new_cat = new_df1.select_dtypes(include="object")
new_num = new_df1.select_dtypes(include="number")
```

```
In [118]: new_cat = encoder.transform(new_cat)
```

```
In [119]: new_cat
```

```
Out[119]:
```

	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	engine
0	13685.540652	13434.368051	13069.674208	13789.064262	12587.501767	20484.76626	18707.867


```
In [120]: new_df1 = pd.concat([new_num, new_cat], axis=1)
new_df1
```

```
Out[120]:
```

	symboling	wheelbase	boreratio	stroke	compressionratio	peakrpm	highwaympg	CarName
0	3	98.0	4.51	4	8.0	5000	29	13685.5406

```
In [121]: new_df1 =pd.DataFrame(scaler.transform(new_df1), columns=new_df1.columns)
new_df1
```

```
Out[121]:
```

	symboling	wheelbase	boreratio	stroke	compressionratio	peakrpm	highwaympg	CarName
0	1.0	0.393103	1.407143	0.919048	0.0625	0.346939	0.342105	0.264484

```
In [122]: pred_new = rfc1.predict(new_df1)
pred_new
```

```
Out[122]: array([19505.56666667])
```

Conclusion

In conclusion, our car price prediction journey involved a comprehensive and systematic process that encompassed various essential steps. We began with Exploratory Data Analysis (EDA), gaining valuable insights into the dataset's characteristics and distributions. This allowed us to better understand the relationships between different features and the target variable, setting the stage for effective modeling.

Following EDA, we carefully partitioned our dataset into training and test sets, ensuring an unbiased evaluation of model performance. Encoding categorical variables and applying feature scaling were crucial steps that facilitated optimal model training by transforming and normalizing the data appropriately.

Hyperparameter tuning played a pivotal role in fine-tuning the models' configurations. This process allowed us to identify the best set of hyperparameters that maximized predictive accuracy and generalization. We meticulously searched through various parameter combinations to strike the right balance between bias and variance.

Feature selection, a key aspect of model refinement, involved choosing the most relevant and informative features to enhance model efficiency and interpretability. By excluding noise and redundant information, we aimed to improve the model's ability to generalize to new, unseen data.

Throughout this journey, the RandomForestRegressor emerged as the most promising model for car price prediction. Its robust performance was a result of its adaptability to complex relationships, demonstrated by its low Mean Absolute Error (MAE), which is a 1178.684498102981, high training score is (97%) 0.9789761134166135, and impressive test score is (94%) 0.942488779940313. This model's ability to capture intricate patterns and its consistent performance across various evaluation metrics make it a strong contender for accurate car price predictions.